




Docker Networking in Production at VISA

Sasi Kannappan

Chief Systems Architect, Visa Inc.

Mark Church

Solutions Architect, Docker Inc.




An abridged history of networking ...



Enter containers ...

- 100s or 1000s of containers per host
- Containers that exist for minutes or months
- Microservices distributed across many more hosts (>>> E-W traffic)

... but this is worse.



Docker Networking

The Container Network Model (CNM)

**Put Users and
Applications First**

.....

Network policies
defined in terms of
applications

Plugin API Design

.....

Batteries included
but removable



Docker Networking

The Container Network Model (CNM)

Networks as a first-class citizen in Docker

Scalable and secure control plane

Pluggable network stack

Support across OS ecosystem

Easy multi-host networking




What is now possible?

docker-stack.yml

```
services:
  frontend:
    networks:
      - ingress
      - backend
    ports:
      - 5000
      - 7000
  database:
    networks:
      - backend

networks:
  ingress:
    driver: overlay
  backend:
    driver: overlay
    driver_opts:
      encrypted : true
```



Docker Networking in Production at VISA: An Evolutionary Story

Sasi Kannappan



History of Docker at Visa

History

Started looking at Docker in late 2015

First containerized application in production in late 2016

Primary Application Environment


- A customer-facing financial transaction platform
- Now in production for 6 months
- Comprised of ~100 containers and ability to scale to ~800
- 2 production, 2 sandbox clusters, across 2 regions

The Vision


Goals

- Migration to Microservices Arch
- Dynamic scalability
- Operational simplicity
- Load balancer-less

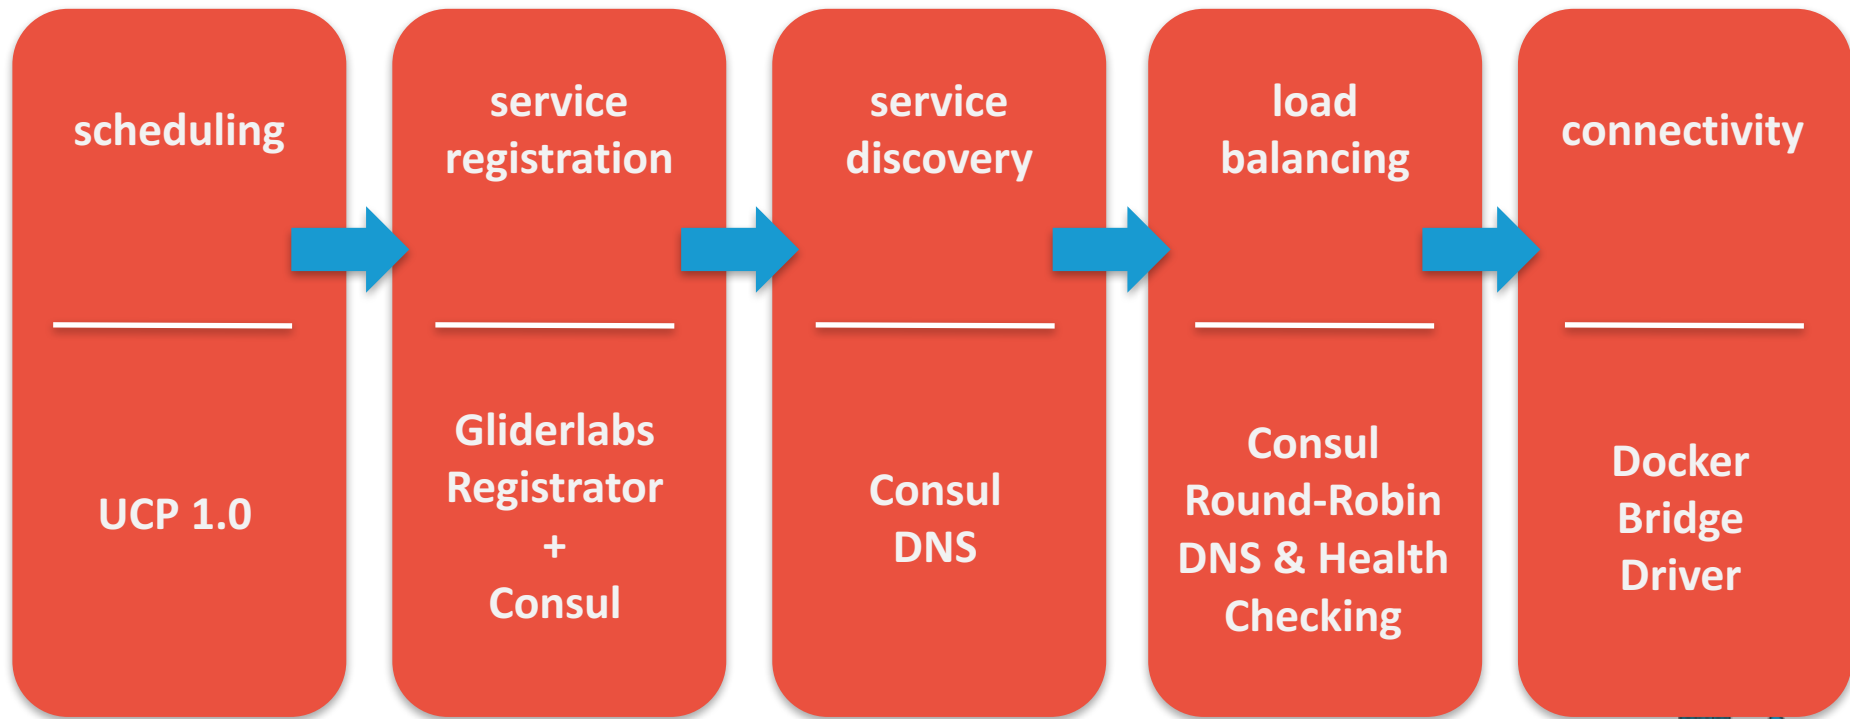
From this ...



to this ...



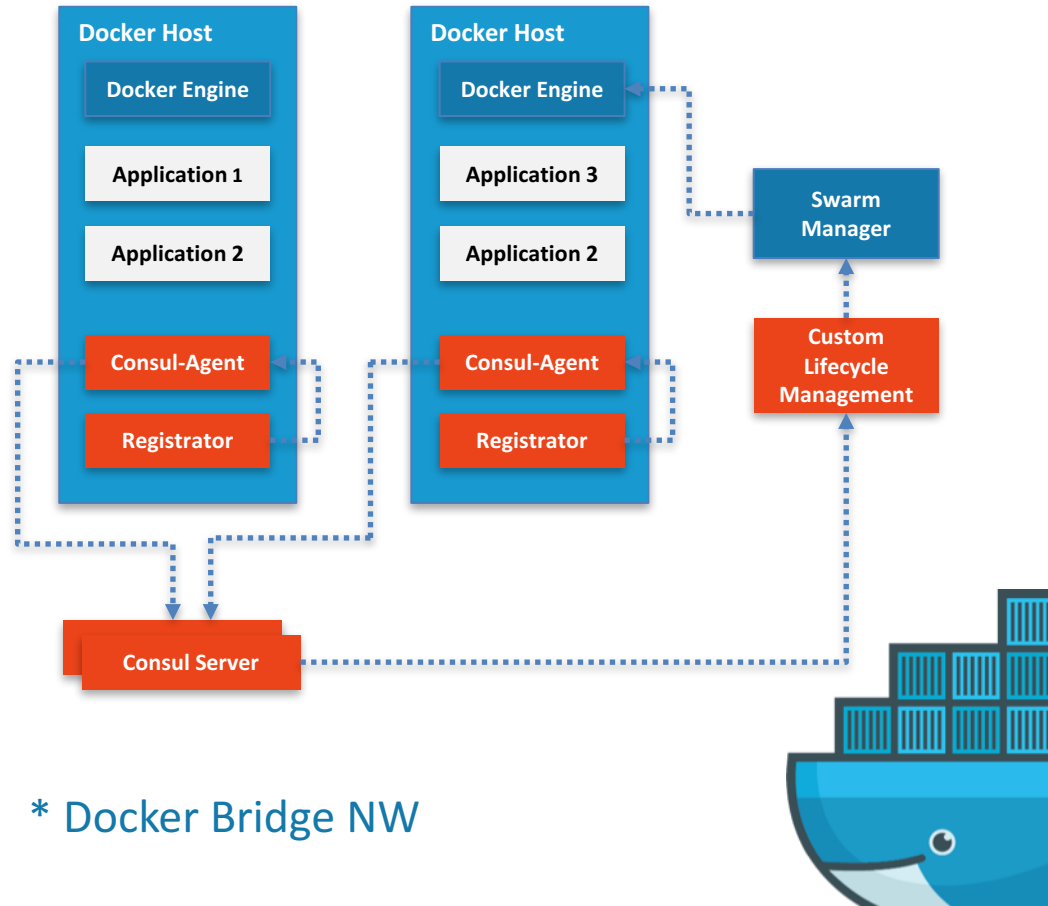
1st Gen Container Networking



1st Gen Docker Architecture

Service Registration

- Services and their location (ip:port) are registered centrally
- 1st gen: handled by registrator and consul

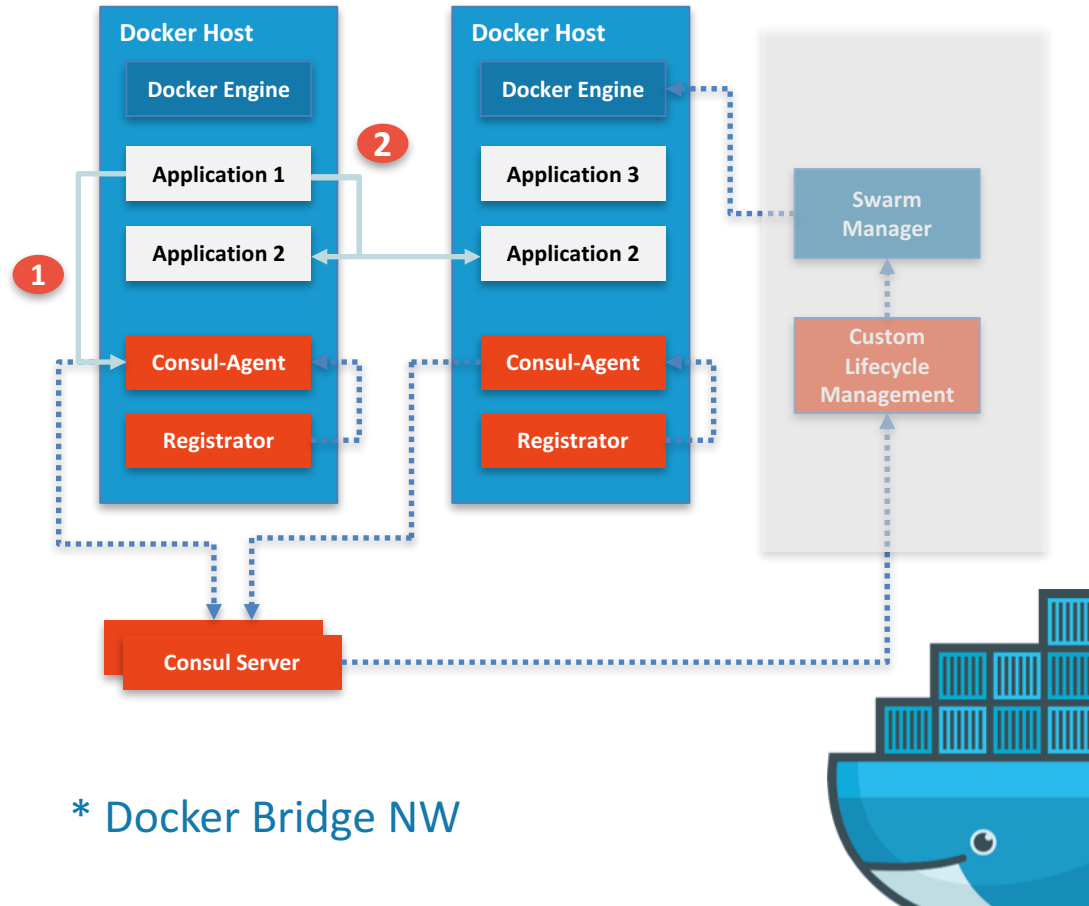


* Docker Bridge NW

1st Gen Docker Architecture

Service Discovery

- The ability for services to find each other
- Designed using Consul capabilities




* Docker Bridge NW

Problems with 1st Gen Architecture

- Complexity
 - Many components to manage
 - Maintaining HA for all components
 - Many component integrations to manage
 - Difficulty in troubleshooting
- Maintainability
 - Custom glue-code to manage

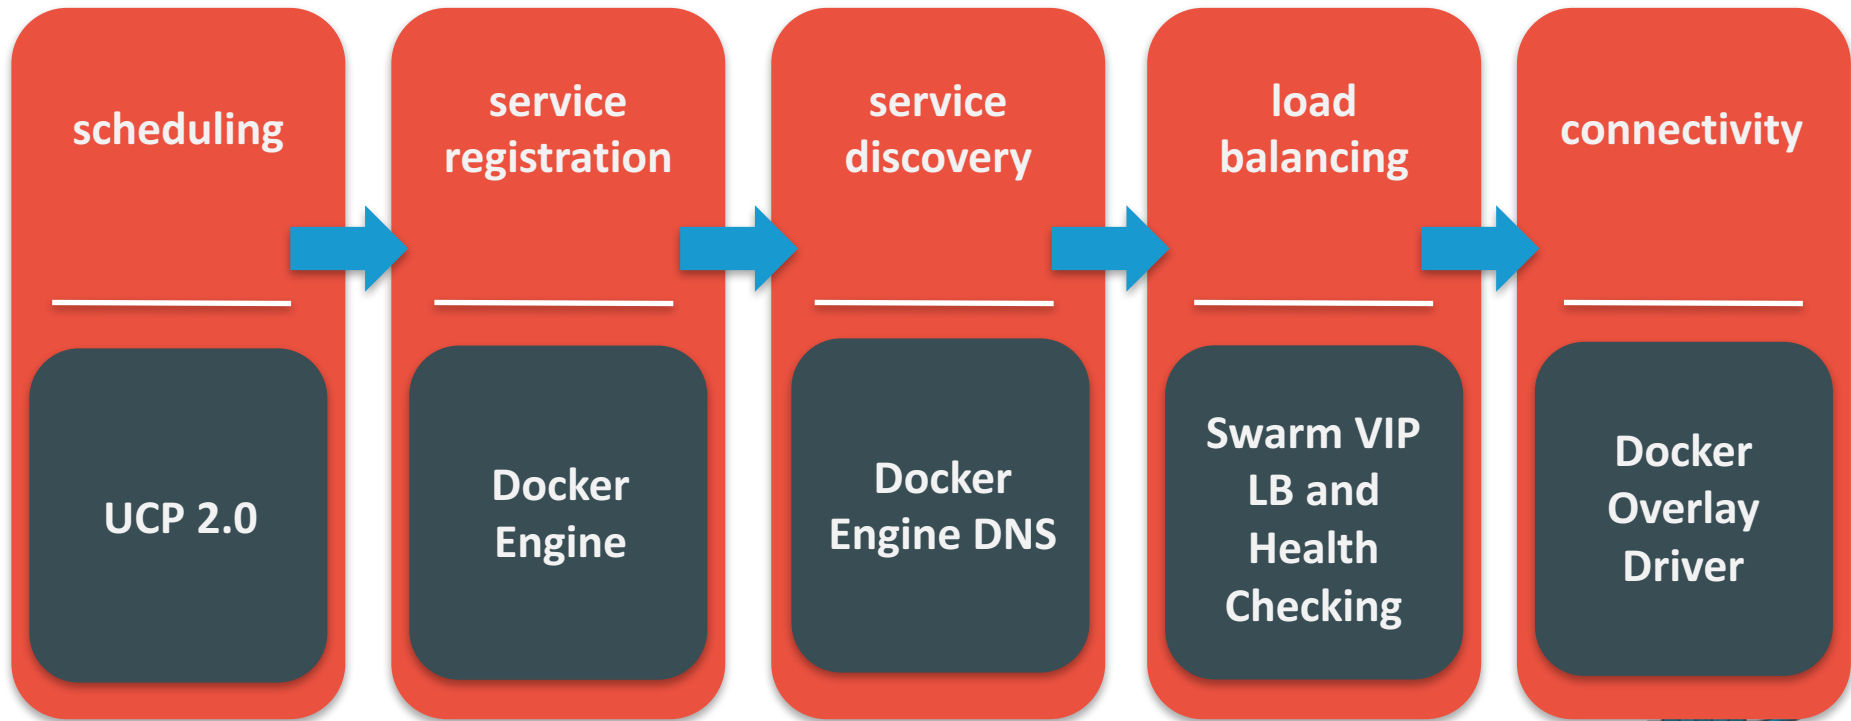
This is the goal ...



but, we replaced
LB with this ...



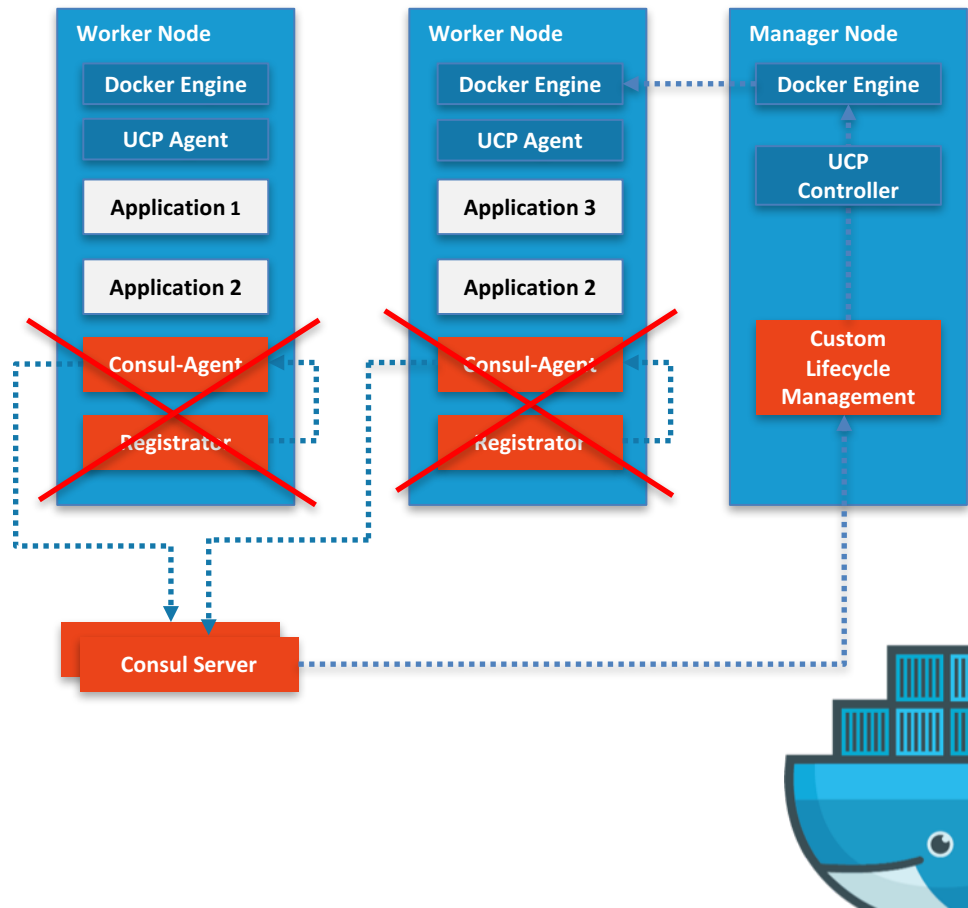
2nd Gen Container Networking



Service Registration & Service Discovery

- Service Definitions in Docker (Built-in registry)
- Docker Overlay NW
- VIP for Services

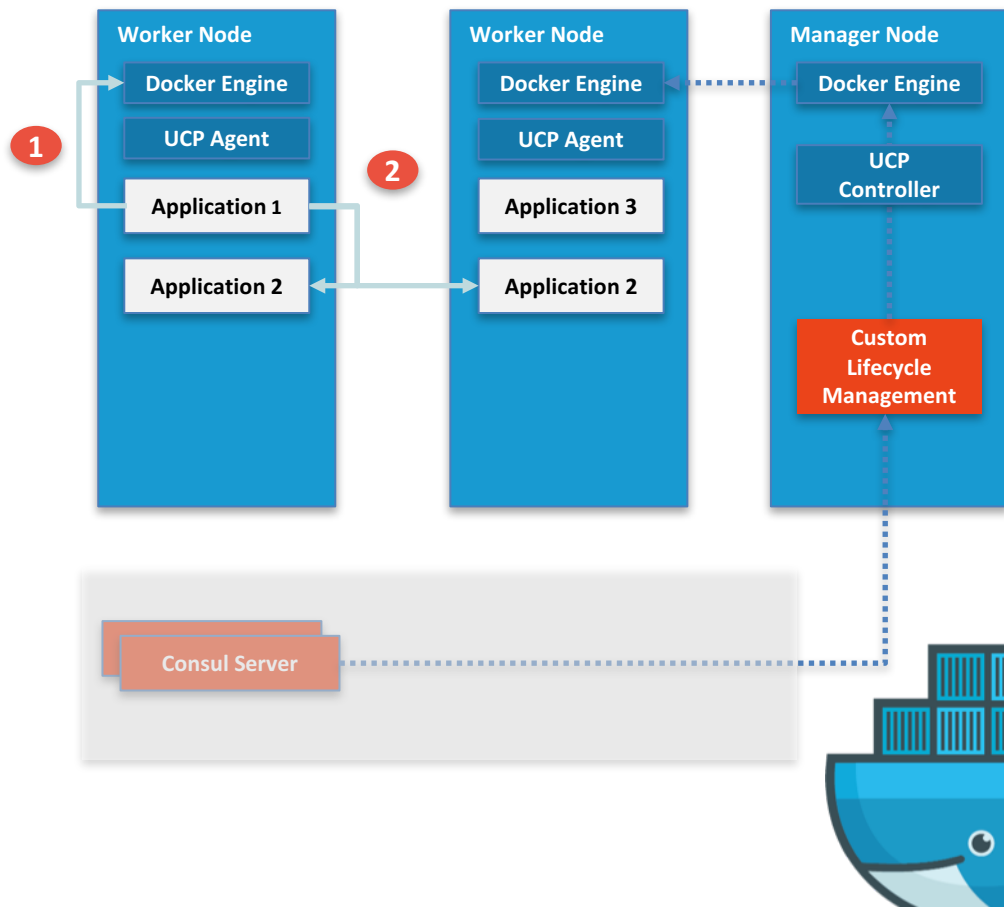
2nd Gen Docker Architecture



Load Balancing & Lifecycle Management

- Transparent to Applications
- Use VIPs
 - application1 (LB to Application2 containers on Nodes 1 & 2)
- [http\(s\)://application1/](http(s)://application1/)

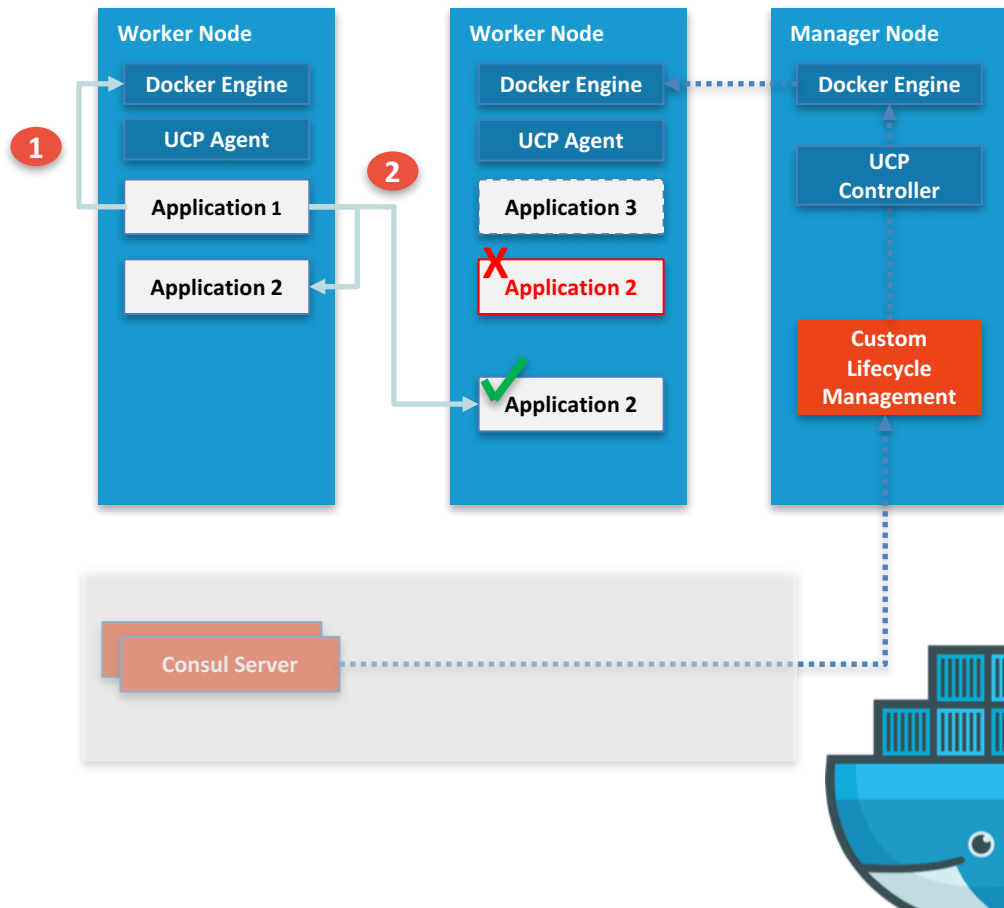
2nd Gen Docker Architecture



Load Balancing & Lifecycle Management

- Integrated Health check
- Self Heal Container Instances

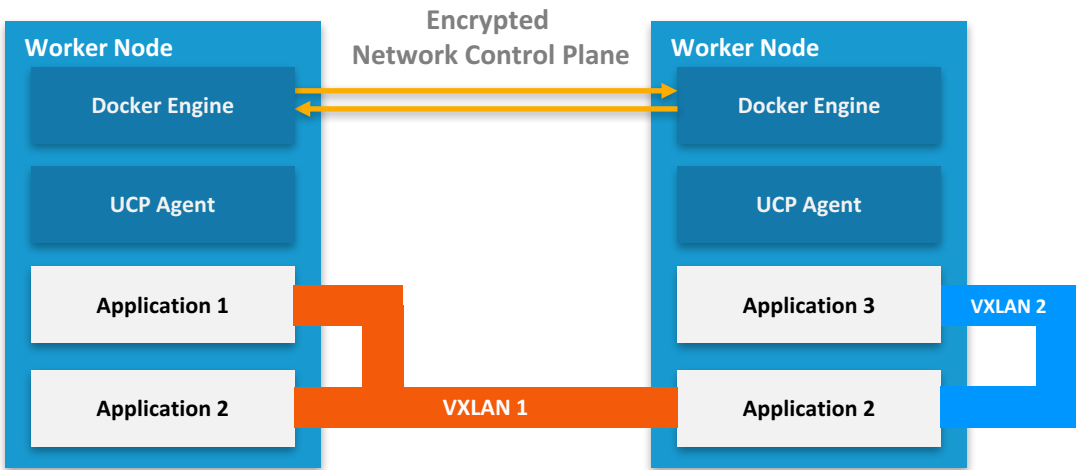
2nd Gen Docker Architecture



Connectivity

- VXLAN Overlay
- Transparent Overlay Encryption

2nd Gen Docker Architecture



Summary

- Less complexity in design -> easier to troubleshoot & better visibility
- Easier to maintain -> less cycles in upgrades and in
- Less custom code & fewer integration points
- Enhanced features (VIP load balancing and lifecycle mgmt)

DEMO TIME!

Thank you!

@docker #dockercon

